

 device: 0, name: Tesla P100-PCIE-16GB, pci bus id:
0000:00:04.0, compute capability: 6.0
```

```
Found 17572 files belonging to 38 classes.
```

```
model = keras.Sequential()
```

```
model.add(keras.layers.Conv2D(32,
(3,3),activation="relu",padding="same",input_shape=(256,256,3)))
```

```

model.add(keras.layers.Conv2D(32,
(3,3),activation="relu",padding="same"))
model.add(keras.layers.MaxPooling2D(3,3))

model.add(keras.layers.Conv2D(64,
(3,3),activation="relu",padding="same"))
model.add(keras.layers.Conv2D(64,
(3,3),activation="relu",padding="same"))
model.add(keras.layers.MaxPooling2D(3,3))

model.add(keras.layers.Conv2D(128,
(3,3),activation="relu",padding="same"))
model.add(keras.layers.Conv2D(128,
(3,3),activation="relu",padding="same"))
model.add(keras.layers.MaxPooling2D(3,3))

model.add(keras.layers.Conv2D(256,
(3,3),activation="relu",padding="same"))
model.add(keras.layers.Conv2D(256,
(3,3),activation="relu",padding="same"))

model.add(keras.layers.Conv2D(512,
(5,5),activation="relu",padding="same"))
model.add(keras.layers.Conv2D(512,
(5,5),activation="relu",padding="same"))

model.add(keras.layers.Flatten())

model.add(keras.layers.Dense(1568,activation="relu"))
model.add(keras.layers.Dropout(0.5))

model.add(keras.layers.Dense(38,activation="softmax"))

opt = keras.optimizers.Adam(learning_rate=0.0001)
model.compile(optimizer=opt,loss="sparse_categorical_crossentropy",metrics=['accuracy'])
model.summary()

```

Model: "sequential"

| Layer (type)                 | Output Shape         | Param # |
|------------------------------|----------------------|---------|
| =====                        |                      |         |
| conv2d (Conv2D)              | (None, 256, 256, 32) | 896     |
| conv2d_1 (Conv2D)            | (None, 256, 256, 32) | 9248    |
| max_pooling2d (MaxPooling2D) | (None, 85, 85, 32)   | 0       |
| conv2d_2 (Conv2D)            | (None, 85, 85, 64)   | 18496   |

|                                |                     |          |
|--------------------------------|---------------------|----------|
| conv2d_3 (Conv2D)              | (None, 85, 85, 64)  | 36928    |
| max_pooling2d_1 (MaxPooling2D) | (None, 28, 28, 64)  | 0        |
| conv2d_4 (Conv2D)              | (None, 28, 28, 128) | 73856    |
| conv2d_5 (Conv2D)              | (None, 28, 28, 128) | 147584   |
| max_pooling2d_2 (MaxPooling2D) | (None, 9, 9, 128)   | 0        |
| conv2d_6 (Conv2D)              | (None, 9, 9, 256)   | 295168   |
| conv2d_7 (Conv2D)              | (None, 9, 9, 256)   | 590080   |
| conv2d_8 (Conv2D)              | (None, 9, 9, 512)   | 3277312  |
| conv2d_9 (Conv2D)              | (None, 9, 9, 512)   | 6554112  |
| flatten (Flatten)              | (None, 41472)       | 0        |
| dense (Dense)                  | (None, 1568)        | 65029664 |
| dropout (Dropout)              | (None, 1568)        | 0        |
| dense_1 (Dense)                | (None, 38)          | 59622    |
| =====                          |                     |          |
| Total params: 76,092,966       |                     |          |
| Trainable params: 76,092,966   |                     |          |
| Non-trainable params: 0        |                     |          |

```

ep = 10
history = model.fit_generator(train_gen,
 validation_data=test_gen,
 epochs = ep)

```

Epoch 1/10

```

2022-06-02 09:49:57.725110: I
tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:185] None of
the MLIR Optimization Passes are enabled (registered 2)
2022-06-02 09:50:00.285178: I
tensorflow/stream_executor/cuda/cuda_dnn.cc:369] Loaded cuDNN version
8005

```

```

2197/2197 [=====] - 316s 140ms/step - loss:
1.6022 - accuracy: 0.5339 - val_loss: 0.5197 - val_accuracy: 0.8419

```

Epoch 2/10

```


2197/2197 [=====] - 195s 89ms/step - loss:
0.4302 - accuracy: 0.8633 - val_loss: 0.2852 - val_accuracy: 0.9066

```

```
Epoch 3/10
2197/2197 [=====] - 194s 88ms/step - loss:
0.2553 - accuracy: 0.9166 - val_loss: 0.2539 - val_accuracy: 0.9214
Epoch 4/10
2197/2197 [=====] - 195s 88ms/step - loss:
0.1754 - accuracy: 0.9418 - val_loss: 0.2177 - val_accuracy: 0.9314
Epoch 5/10
2197/2197 [=====] - 193s 88ms/step - loss:
0.1299 - accuracy: 0.9573 - val_loss: 0.1694 - val_accuracy: 0.9451
Epoch 6/10
2197/2197 [=====] - 194s 88ms/step - loss:
0.1014 - accuracy: 0.9661 - val_loss: 0.1315 - val_accuracy: 0.9583
Epoch 7/10
2197/2197 [=====] - 194s 88ms/step - loss:
0.0823 - accuracy: 0.9719 - val_loss: 0.1392 - val_accuracy: 0.9585
Epoch 8/10
2197/2197 [=====] - 204s 93ms/step - loss:
0.0691 - accuracy: 0.9769 - val_loss: 0.1477 - val_accuracy: 0.9547
Epoch 9/10
2197/2197 [=====] - 194s 88ms/step - loss:
0.0623 - accuracy: 0.9796 - val_loss: 0.1131 - val_accuracy: 0.9668
Epoch 10/10
2197/2197 [=====] - 194s 88ms/step - loss:
0.0532 - accuracy: 0.9829 - val_loss: 0.1166 - val_accuracy: 0.9677
```

```
plt.figure(figsize = (20,5))
plt.subplot(1,2,1)
plt.title("Train and Validation Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.plot(history.history['loss'],label="Train Loss")
plt.plot(history.history['val_loss'], label="Validation Loss")
plt.xlim(0, 10)
plt.ylim(0.0,1.0)
plt.legend()

plt.subplot(1,2,2)
plt.title("Train and Validation Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.plot(history.history['accuracy'], label="Train Accuracy")
plt.plot(history.history['val_accuracy'], label="Validation Accuracy")
plt.xlim(0, 9.25)
plt.ylim(0.75,1.0)
plt.legend()
plt.tight_layout()
```



```

labels = []
predictions = []
for x,y in test_gen:
 labels.append(list(y.numpy()))
 predictions.append(tf.argmax(model.predict(x),1).numpy())

predictions = list(itertools.chain.from_iterable(predictions))
labels = list(itertools.chain.from_iterable(labels))

print("Train Accuracy : {:.2f} %".format(history.history['accuracy']
[-1]*100))
print("Test Accuracy : {:.2f} %".format(accuracy_score(labels,
predictions) * 100))
print("Precision Score : {:.2f} %".format(precision_score(labels,
predictions, average='micro') * 100))
print("Recall Score : {:.2f} %".format(recall_score(labels,
predictions, average='micro') * 100))

Train Accuracy : 98.29 %
Test Accuracy : 96.77 %
Precision Score : 96.77 %
Recall Score : 96.77 %

plt.figure(figsize= (20,5))
cm = confusion_matrix(labels, predictions)
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
 display_labels=list(range(1,39)))
fig, ax = plt.subplots(figsize=(15,15))
disp.plot(ax=ax,colorbar= False,cmap = 'YlGnBu')
plt.title("Confusion Matrix")
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()

```

<Figure size 1440x360 with 0 Axes>

